

Topology as a Computational Basis for ICL in First-Order CRNs

A concise account of the theory, experiments, audit, and current evidence

Topology worktree: `topology`

Evidence source commit: `556a4d1`

Markov exact-control update: `2c4cba6`

Multi-base handoff update: `d849378`

May 13, 2026

Summary. For first-order CRNs with exponential input-dependent rates, topology is not just a container for trainable parameters. It determines the rooted-spanning-tree sums through which edge projections enter the steady state. The computational basis is therefore not the set of edge vectors $\{K_e\}$, but the set of tree-sum vectors

$$\Theta_T = \sum_{e \in T} K_e, \quad T \in \mathcal{T}_r(G), \quad r \in V.$$

The experiments support the scoped statement that, in the tested fixed-count regimes, topology-associated structural and functional variables explain residual novel-class ICL variation that raw trainable count does not explain. The exact-control phase sharpens this: repaired no-bias γ_{ICL}^* passes analytic toy gates but is not yet predictive, prospective tree-difference overlap did not survive the first exact causal control, and normal-fan/tree-count geometry is now the best weak pre-training signal. The current handoff has built the next multi-base exact-degree library to separate rooted-tree abundance from task-aligned normal-fan geometry. There is still no universal scalar topology law.

Scope. Everything below is about first-order CRNs with exponential rate encoding. The matrix-tree theorem makes this case exact. Autocatalytic and WTA CRNs need separate topology theories.

[1] The starting question

The original CRN-ICL paper showed that a chemical network can do in-context learning without transformer attention. The input is a context/query vector

$$z = (z_1, \dots, z_{N_c}, z_q), \quad z_i \in \mathbb{R}^D,$$

and the target is the context position whose item matches the query. The true metric is novel-class ICL accuracy, not training accuracy or ordinary validation accuracy.

The paper’s mechanism was subspace projection. For each context position i , define a comparison region

$$\mathcal{M}_i = \{z : z_i \sim z_q\}.$$

ICL succeeds when inputs from different comparison branches produce steady states that the decoder separates.

The original degree-count story was broadly right:

$$n_{\text{req}} = 2N_c(N_c + 1)D, \quad n_g = N_n(N_n - 1)(N_c + 1)D$$

for fully connected first-order CRNs. But this count does not say where the degrees of freedom sit. The topology question was therefore

At fixed count, does the reaction/input topology change ICL?

[2] The exact first-order object

Let $G = (V, E)$ be a strongly connected directed first-order reaction graph. For edge $e : u \rightarrow v$, use exponential rate encoding

$$k_e(z) = \exp(b_e + K_e^\top z).$$

The first-order steady state is controlled exactly by the matrix-tree theorem:

$$\bar{C}_r(z) = \frac{\tau_r(z)}{\sum_s \tau_s(z)}, \quad \tau_r(z) = \sum_{T \in \mathcal{T}_r(G)} \prod_{e \in T} k_e(z).$$

Substitute the exponential form:

$$\prod_{e \in T} k_e(z) = \exp \left(\sum_{e \in T} b_e + \left(\sum_{e \in T} K_e \right)^\top z \right).$$

Set

$$\beta_T = \sum_{e \in T} b_e, \quad \Theta_T = \sum_{e \in T} K_e.$$

Then

$$\tau_r(z) = \sum_{T \in \mathcal{T}_r(G)} \exp(\beta_T + \Theta_T^\top z).$$

This is the key point. A first-order CRN does not read out isolated edge projections $K_e^\top z$. It reads out log-sum-exp combinations of rooted-tree projections $\Theta_T^\top z$.

[3] What topology changes


Stack tree incidences as row vectors $s_T \in \{0, 1\}^m$. Then

$$\Theta_T = K^\top s_T.$$

Since concentrations are normalized, common shifts of all tree scores cancel. The more relevant structural object is the tree-difference matrix

$$D_G = \begin{bmatrix} s_{T_1} - s_{T_0} \\ s_{T_2} - s_{T_0} \\ \vdots \end{bmatrix}.$$

physical graph $G = (V, E)$



matrix-tree

computational basis

$$au_r(z) = \sum_{T \in \text{Trees}_r(G)} e^{\beta_T + \Theta_T^\top z}$$

$$\Theta_T = \sum_{e \in T} K_e$$

basis is $\{\Theta_T : T \in \text{Trees}_r(G), r \in V\}$
 not the isolated edge set $\{K_e : e \in E\}$

edges carry learned vectors K_e

Figure 1: Topology changes the projection basis. Edge vectors K_e are learned locally, but the steady state uses rooted tree-sum projections Θ_T .

A first count of relative projection capacity is

$$d_{\text{rel}}(G) = p \text{rank}(D_G), \quad p = (N_c + 1)D.$$

With an input mask $\Omega_{e\alpha}$, the capacity must be input-aware:

$$d_{\text{rel}}(G, \Omega) = \sum_{\alpha=1}^p \text{rank}(D_G \text{diag}(\Omega_{\cdot\alpha})).$$

But this is still only a rank proxy. ICL asks a sharper question:

Can the allowed tree-sum directions separate the comparison branches?

The capacity object should be closer to

$$\gamma^*(G, \Omega) = \max_{K, B} \min_{z \in \mathcal{B}} \left[q_{\ell^*(z)}(z) - \max_{\ell \neq \ell^*(z)} q_{\ell}(z) \right]$$

under norm constraints. In the tropical limit,

$$\log \tau_r(z) \approx \max_{T \in \mathcal{T}_r(G)} (\beta_T + \Theta_T^\top z),$$

so the theory becomes a rooted-tree-polytope normal-fan coverage problem.

The hierarchy that emerged is therefore

raw count < relative tree rank < masked tree geometry

< branch-margin/tree-polytope geometry < trained organization < causal dependence.

[4] What was built and audited

The worktree now contains code and reports for arbitrary first-order directed topologies, input masks, tree metrics, branch-capacity probes, post-training tree diagnostics, ablations, motif extraction, and statistic-preserving causal scrambles.

The audit report `ICL/results/next_phase_stats/topology_goal_completion_audit.md` maps the goal to concrete artifacts. It records that the verifier passed for committed artifacts, that hard regimes have full result/model/mechanism/causal follow-ups, and that all headline claims are scoped to `test_novel_classes`. The audit also marks the branch-margin theory as partial by design: proxies exist, but there is not yet a final theorem.

The main source reports used here are:

<code>topology_icl_complete_explanation.md</code>	full narrative explanation
<code>next_phase_evidence_report.md/json</code>	fixed-count, motif, causal summaries
<code>mechanism_isolation_evidence.csv</code>	controlled mechanism contrasts
<code>stat_preserving_causal_stratified</code>	stratified causal scrambles
<code>degree_rewire_normal_fan...</code>	exact-degree normal-fan pilot

[5] Fixed-count structural evidence

The original controlled regime fixed

$$N_n = 6, \quad m = 20, \quad \text{input-coupled count} = 200.$$

It used three physical backbones, 16 selected masks/topologies per backbone, five training seeds each, 240 total runs, and 48 topology/mask groups.

At group level, with target equal to mean novel-class ICL accuracy, the leave-one-out values were

model	group LOO R^2	interpretation
raw count	-0.043	no explanatory power at fixed count
raw count + d_{rel}	0.145	relative rank helps
masked tree geometry	0.189	input-aware geometry helps
tree geometry	0.409	strongest structural summary here

The corresponding hard $N_n = 5, m = 12, N_c = 3, D = 2$ regime showed a different but still positive pattern: raw count and d_{rel} alone were both -0.190, tree geometry was 0.324, and masked tree geometry was 0.447.

The safe conclusion is not that d_{rel} is the law. It is that topology-associated variables explain residual ICL variation that raw count cannot explain in the tested regimes.

[6] Why input masks matter

Physical topology and input-encoding topology are distinct. The physical graph fixes the rooted trees:

$$G_{\text{phys}} \mapsto \mathcal{T}_r(G).$$




Figure 2: Structural predictors in two fixed-count regimes. Raw count is uninformative once fixed. Tree geometry and masked tree geometry explain residual group-level variation, but the best predictor depends on regime.

The input mask fixes which coordinates can enter which edge vectors:

$$\Omega_{e\alpha} = 0 \quad \Rightarrow \quad K_{e\alpha} = 0.$$

The real capacity object is therefore $\gamma^*(G, \Omega)$, not only $\gamma^*(G)$.

The mechanism-isolation file shows fixed-graph, fixed-count, fixed- d_{rel} cases where masks still change ICL. On the fixed random physical graph, changing edge-load structure gave +6.96 mean ICL points, while changing coordinate-load heterogeneity gave -16.48 mean ICL points. On the fixed cycle graph, the corresponding deltas were -6.32 and +2.12 points; on the fixed hub graph they were +3.44 and -1.52 points. The sign is not universal, but the effect exists.

[7] Post-training mechanism evidence

Pre-training topology asks what the graph could use. Post-training diagnostics ask what the trained model did use. The strongest diagnostic class was projection/root/tree organization around comparison branches. Projection-alignment summaries gave

$$\begin{aligned} R_{LOO}^2 &= 0.740 && \text{at run level,} \\ R_{LOO}^2 &= 0.809 && \text{at topology-mean level,} \\ R_{LOO}^2 &= 0.599 && \text{for topology-best seed.} \end{aligned}$$

This means high-ICL models are not merely high-count systems. They organize the available tree-geometry so that comparison branches produce decodable steady states. The active-tree and




Figure 3: Controlled contrasts. Some coarse structural metrics are not monotone drivers, and fixed-graph input masks can still change ICL. This is why d_{rel} alone is incomplete.




Figure 4: Post-training projection alignment is much more explanatory than coarse pre-training count. This is mechanism evidence, not an independent capacity theorem.

margin diagnostics are close to the learned mechanism, so they should be read as mechanistic explanations rather than clean pre-training predictors.

[8] Causal scrambles

The causal interventions are the strongest evidence that the learned alignment matters. Two statistic-preserving scrambles were central:

branch-alignment scramble	preserves graph, root tree counts, d_{rel} , and input density
projection scramble	also preserves input support and per-edge projection row norms

They break branch/projection/tree organization while preserving many coarse statistics.

In the stratified hard-regime pilot, high-ICL runs dropped by about -96 points under branch-alignment scrambling and about -64 to -68 points under projection scrambling. Low and mid

runs also collapsed substantially.




Figure 5: Statistic-preserving scrambles in low, mid, and high baseline buckets. Destroying branch/projection/tree alignment sharply reduces novel-class ICL.

This pushes the evidence beyond correlation. It says the trained CRNs rely on the organized relation between input branches, learned edge projections, and root/tree structure.

[9] Hard regimes and the normal-fan pilot

Three hard regimes were added: $n4_m6_N3_D2$, $n5_m8_N3_D2$, and $n5_m12_N3_D2$, each with 60 trained runs plus mechanism and causal follow-ups. The results are mixed in the useful sense. Topology signal depends on regime.

In $n5_m12_N3_D2$, masked tree geometry had LOO $R^2 = 0.447$ and positive held-out-family performance. In the sparser $n5_m8_N3_D2$ regime, within-sample bootstrap improvements were more positive than held-out-family prediction. This suggests that topology may have sparse, intermediate, and redundant regimes rather than one monotone scalar rule.

The exact-degree normal-fan pilot fixed

$$N_n = 5, \quad m = 12, \quad N_c = 3, \quad D = 2, \quad \text{in/out degree sequence} = [3, 2, 2, 3, 2], \quad d_{\text{rel}} = 88.$$

It varied tree-polytope/normal-fan diagnostics through directed degree-preserving rewires. Only four topology groups were trained, so this is not yet a statistical result. It is a constructive design showing that topology can vary beyond degree sequence and d_{rel} .

The descriptive $n = 4$ correlations were $r = 0.869$ for active-tree count and $r = 0.691$ for branch-tree NMI with mean ICL. These numbers are promising design feedback, not proof.




Figure 6: Exact-degree pilot. All points have the same degree sequence and $d_{\text{rel}} = 88$. Active tree count and branch-tree NMI vary across rewires, and mean ICL varies with them in this small $n = 4$ pilot.

[10] Motifs and matched controls

Essential physical subgraphs retrained better than sparse input masks under the extraction protocol. But matched controls softened the stronger interpretation. Degree-rewired and random strongly connected controls often matched or exceeded the extracted motifs.

The corrected interpretation is:

small physical subgraphs can support ICL, but the extracted motifs are not uniquely optimal.

This is compatible with the tree-sum view: many different graphs can provide usable branch-covering tree geometry.

[11] The gamma-star attempt

The project implemented γ^* -style capacity probes because d_{rel} is too coarse. This was the right theoretical direction, but the first optimized proxy did not outperform tree geometry. In the hard $n5_m12_N3_D2$ regression, optimized γ capacity had LOO $R^2 = -0.564$, while tree geometry had 0.324 and masked tree geometry had 0.447.

The next audit repaired the analytic toy gate. The no-bias probe now reproduces the paper’s small cases: two species with both branches fails, two species with the max branch passes, and the complete directed three-species both-branch construction passes. That repair validates the lower-tail object as a diagnostic. It does not make γ_{ICL}^* a predictor: in the fixed m20 reanalysis, tree-difference multiplicity reached mean-ICL LOO $R^2 = 0.435$, while the best repaired gamma model reached only 0.078.




Figure 7: Extracted physical motifs support ICL, but matched controls are competitive. The evidence supports a family of feasible physical substrates, not one unique motif.

The lesson is sharper than simple failure. ICL is a worst-branch problem:

$$\max_{K,B} \min_{b \in \mathcal{B}} \text{margin}_b(K, B; G, \Omega).$$

A good capacity theory should target low-percentile or worst-branch separation under norm constraints, but the current repaired γ diagnostic still misses important topology/mask variables in trained regimes.

[12] The Markov-ICL extension

The next question was not “run a bigger sweep.” It was:

Can the first-order topology story be sharpened using Markov-process expressivity theory?

The requested ingredients were input multiplicity, monotonicity limits, coefficient constraints, non-equilibrium driving, branch sharpness, structural compatibility, and a clean separation between expressivity and trainability.

The exact representation from Section 2 is already the bridge. For first-order CRNs,

$$\tau_r(z) = \sum_{T \in \mathcal{T}_r(G)} \exp(\beta_T + \Theta_T^\top z), \quad \Theta_T = \sum_{e \in T} K_e.$$

Thus the Markov-expressivity objects must be defined on tree-sum features, not isolated edges.

The input-encoding multiplicity is

$$M_\alpha = \sum_e \Omega_{e\alpha}.$$

For ICL, the important refinements are branch-specific:

$$M_{i,d}, \quad M_{q,d}, \quad \sum_e \Omega_{e,i,d} \Omega_{e,q,d}, \quad |M_{i,d} - M_{q,d}|.$$

Low or imbalanced comparison-coordinate multiplicity is the monotonicity-risk diagnostic:

$$\text{monoRisk}_\alpha = \mathbf{1}[M_\alpha \leq 1].$$

After training, an effective participation diagnostic is

$$M_\alpha^{\text{eff}} = \frac{(\sum_e |K_{e\alpha}|)^2}{\sum_e K_{e\alpha}^2}.$$

Branch sharpness is distinct from branch coverage. For branch direction u_b ,

$$R_{r,b} = \max_{T \in \mathcal{T}_r(G)} \Theta_T^\top u_b - \min_{T \in \mathcal{T}_r(G)} \Theta_T^\top u_b.$$

This measures whether a rooted tree polytope can create a sharp margin along a comparison direction. A graph can cover every branch and still have weak lower-tail margins.

Coefficient controllability is the other missing piece. Tree-polynomial coefficients are constrained by overlapping spanning trees; they are not freely adjustable. Therefore rank, even masked rank, can overestimate useful capacity. Effective rank, condition number, tree entropy, normal-fan coverage, extremal-tree accessibility, and branch-specific concentration are all relevant.

The theory note produced in this phase is `ICL/markov_icl_expressivity_theory.md`. Its shortest message is:

First-order Markov-ICL expressivity is tree-sum expressivity under an input mask. Input multiplicity and branch sharpness are useful structural variables, but not standalone laws.

[13] Existing-data Markov reanalysis

The reanalysis script `ICL/markov_expressivity_reanalysis.py` consumes committed CSV/JSON artifacts only. It does not launch new broad training. It groups seeds by topology/mask before inference.

Coverage:

dataset	rows	topology/mask groups
fixed m20	240	48
hard $n4_m6$	60	12
hard $n5_m8$	60	12
hard $n5_m12$	60	12
exact-degree pilot	4	4
exact-degree library	32	32

In the fixed m20 data, grouped LOO R^2 for mean novel-class ICL was 0.095 for aggregate multiplicity variables, 0.046 for comparison-multiplicity variables, 0.158 for tree-geometry variables, and 0.145 for the old capacity proxy. In the hard $n5_m12$ data, tree geometry reached 0.437 while aggregate multiplicity and the old capacity proxy were -0.190 .




Figure 8: Existing-data Markov reanalysis. The richer Markov variables are useful diagnostics, but the evidence still does not support a universal scalar expressivity law.

The input-multiplicity control is especially clarifying. In the fixed-count mask data, all mask families have mean aggregate multiplicity $M_{\text{mean}} = 10$. Yet the coordinate-block family has zero minimum comparison overlap, high coordinate-load Gini, and much lower mean ICL: 63.4 points versus roughly 74–77 points for the other families.

The exact-degree normal-fan pilot was also reinterpreted through this lens. It has only four trained groups, but those groups share the intended fixed-degree and fixed- d_{rel} controls. Descriptive correlations with mean ICL were $r = 0.869$ for normal-fan active tree count, $r = 0.691$ for branch-tree NMI, $r = 0.810$ for log tree count, and $r = 0.759$ for effective rank. These are design signals, not statistical proof.




Figure 9: Input multiplicity control from existing data. Average multiplicity alone is not the law; branch comparison overlap and load imbalance are the direct failure risks.

[14] Improved lower-tail capacity

The improved probe is ICL/branch_margin_capacity_v2.py. It targets the lower tail directly:

$$\gamma_{\text{ICL}}^*(G, \Omega; \mathcal{C}) = \max_{\theta} \min_{b \in \mathcal{B}} \text{LCVaR}_{\alpha} [m_{\theta}(z) \mid z \in b],$$

where

$$m_{\theta}(z) = q_{y(z)}(z) - \max_{\ell \neq y(z)} q_{\ell}(z).$$

The implemented variants are exact log-sum-exp tree features, tropical max-over-tree features, and hard-root structural compatibility. The constraints include row-norm bounds for K_e , decoder Frobenius norm bounds, edge-bias bounds, and the input mask Ω .

A small smoke run on a 5-node, 12-edge cycle-chords graph produced finite branch-wise reports for all three variants. The objectives were still negative: -3.001 for exact, -2.800 for tropical, and -2.270 for hard-root. This validated the reporting path and reinforced that average accuracy is not the right target; the worst branch can still fail.




Figure 10: Lower-tail capacity V2 smoke result. The new probe reports the correct object: worst-branch lower-tail margin and branch failures. This smoke run is not an optimized capacity theorem.

The later toy repair passed the analytic gate, but the existing-data reanalysis remained negative for prediction. Repaired no-bias exact, tropical, and hard-root gamma variants had fixed-m20 mean-ICL LOO $R^2 = -0.075, -0.115, \text{ and } -0.150$, respectively. Combining repaired gamma with tree-difference multiplicity degraded the fixed-m20 mean-ICL model from 0.435 to 0.078. Thus γ_{ICL}^* should remain a diagnostic until it predicts held-out exact-control outcomes.

[15] Expressivity, trainability, and thermodynamics

The reanalysis keeps three outcomes separate:

best seed $\approx$ expressivity envelope,
 mean seed $\approx$ trainability/reliability,
 seed variance $\approx$ optimization instability.

Across the 36 hard-regime groups, masked effective rank correlated 0.895 with best seed and 0.897 with mean seed, but only 0.345 with seed standard deviation. Trained margin and tree entropy were also strongly correlated with best/mean seed, but they are post-training mechanism descriptors rather than pre-training capacity metrics.




Figure 11: Expressivity and trainability should not be collapsed into one outcome. Best seed, mean seed, and seed standard deviation respond differently to structural and trained-mechanism variables.

The thermodynamic question is deliberately unresolved. Existing arbitrary directed exponential-rate models may be non-equilibrium, but they were not parameterized as reversible-edge thermodynamic Markov processes. Therefore the report `thermodynamic_force_budget_report.md/json` does not claim entropy production or a force-budget law. It only audits eligibility: among 36 local hard topology groups with topology JSON, the mean reversible-edge fraction is 0.512, with range 0 to 1.

A valid thermodynamic experiment needs a reversible support and a parameterization such as

$$W_{ij} = \exp(E_j - B_{ij} + F_{ij}/2 + \text{input drive}), \quad B_{ij} = B_{ji}, \quad F_{ij} = -F_{ji}, \quad |F_{ij}| \leq F_{\max}.$$

Only then should we sweep $F_{\max}$ and ask whether ICL branch margins increase with allowed non-equilibrium driving.

[16] Post-gamma exact-control phase

The next phase was deliberately gated. It did not launch another broad topology sweep. It first reconciled predictor names, lifted input multiplicity into tree and tree-difference space, repaired the analytic γ_{ICL}^* toy gate, and only then ran prospective exact controls.

The fixed-m20 tree-level multiplicity reanalysis was strong:

outcome	edge multiplicity	tree multiplicity	tree-difference multiplicity
mean novel-class ICL	-0.002	0.403	0.435
best-seed novel-class ICL	0.109	0.245	0.419

This supports the statement that useful input multiplicity lives in rooted-tree and tree-difference space, not merely edge space.

The repaired gamma gate passed the analytic small cases, but the existing-data prediction test was negative. In fixed m20, tree-difference multiplicity reached mean-ICL LOO $R^2 = 0.435$, while the best repaired gamma model reached only 0.078. Repaired exact, tropical, and hard-root gamma alone were all negative.




Figure 12: Repaired no-bias gamma passes analytic toys but does not yet predict trained ICL in existing data. Tree-difference multiplicity remains stronger in fixed m20.

The first prospective tree-difference causal control then fixed one physical graph, input-coupled count, d_{rel} , aggregate multiplicity, edge load, and coordinate-load stratum. It trained 16 mask groups with five seeds each. The intended high-overlap contrast did not improve ICL:

stratum	outcome	high overlap	low overlap	high-low
balanced load	mean ICL	75.360	79.330	-3.970
balanced load	best ICL	84.900	89.700	-4.800
imbalanced load	mean ICL	68.830	69.900	-1.070
imbalanced load	best ICL	74.450	76.250	-1.800

The balanced mean-ICL CI was $[-7.9905, 0.6700]$, and the balanced best-ICL CI was $[-9.000, -0.150]$.

Prospective exact control did not support tree-difference as a standalone knob




Figure 13: Prospective tree-difference exact control. High same-root tree-difference comparison overlap did not causally improve novel-class ICL under the first balanced-load exact-control design.

Grouped LOO tells the same story. For mean ICL, controls alone reached 0.488, tree-difference plus controls reached 0.447, and gamma plus controls reached 0.320. For best ICL, controls alone reached 0.601, tree-level plus controls reached 0.631, tree-difference plus controls reached 0.545, and gamma plus controls reached 0.571.

[17] Exact-degree normal-fan expansion

The exact-degree normal-fan pilot was then scaled. The expanded experiment fixed

$$N_n = 5, \quad m = 12, \quad N_c = 3, \quad D = 2,$$

the exact in-degree and out-degree sequence, full input coupling, $d_{\text{rel}} = 88$, and the full-coupling multiplicity distribution. It trained 32 topology groups with five seeds per group, for 160 runs.

This is now the best pre-training signal, but it is weak. Grouped LOO values were:

model	mean ICL	best ICL	seed std.
active-tree count	0.099	0.091	-0.069
normal-fan pair	0.112	0.110	-0.101
tree count	0.114	0.093	-0.069
gamma exact	-0.127	-0.072	-0.105
gamma + normal fan	0.059	0.117	-0.126




Figure 14: Scaled exact-degree normal-fan experiment. Active-tree/tree-count variables are the best current pre-training signal under exact controls, but the effect is small and does not justify a universal scalar law.

Correlations were clearer than LOO prediction: active-tree count correlated about 0.446 with mean ICL and 0.441 with best ICL; log rooted-tree count correlated about 0.460 with mean ICL and 0.439 with best ICL. This suggests that the next experiment must separate total rooted-tree count from genuine normal-fan branch geometry across multiple base graphs or degree sequences.

[18] Updated mechanism evidence

The exact-control phase also added post-training mechanism and causal intervention checks for the prospective-control models. Mechanism metrics were collected for all 80 runs. Four selected high-performing models were then subjected to statistic-preserving scrambles and selected edge ablations.

The causal drops were large:

intervention	mean accuracy change
context-block shuffle	−71.458
stat-preserving projection scramble	−59.688
stat-preserving branch-alignment scramble	−57.778
decoder-root permutation	−53.125

Selected edge ablations also showed dependence on input-coupled and physical edges: mean max input-ablation loss was 21.146 points, and mean max physical-edge ablation loss was 26.563 points.




Figure 15: Mechanism follow-up after the prospective exact-control phase. Trained successful models still depend on branch/projection organization and ablation-sensitive edges.

This is an important separation. The prospective pre-training tree-difference overlap contrast was negative or inconclusive, but trained models that do work still causally depend on branch/projection organization. The mechanism claim is stronger than the current pre-training selector claim.

[19] Current synthesis and next step

The current claim separation is:

category	current status
Expressivity	First-order tree-sum basis is exact. Normal-fan active-tree/tree-count variables weakly predict best-seed ICL under exact controls.
Trainability	Mean-seed ICL follows the same weak normal-fan/tree-count signal, but seed variance is not explained well.
Mechanism	Trained models rely on branch/projection/tree organization; statistic-preserving scrambles collapse selected trained models.
Causal evidence	The first prospective tree-difference causal control did not support tree-difference overlap as a standalone causal knob.
Thermodynamics	Untested. No $F_{\max}$ claim should be made until a reversible-support parameterization is implemented.

The concrete next task is no longer to prove that topology matters broadly. As of the latest handoff, it is:

diagnose why the fixed-m20 tree-difference signal failed under prospective exact control, and test whether normal-fan active-tree/tree-count geometry survives multi-base exact controls after separating it from total rooted-tree count.

The next model should therefore:

1. treat same-root tree-difference overlap as a failed standalone causal knob, not as the main selector;
2. use cross-root and decoder-aware contrast metrics as diagnostics, not as proof;
3. consume the new multi-base exact-degree library that separates rooted-tree count from normal-fan geometry;
4. keep repaired γ_{ICL}^* as a diagnostic rather than a selector;
5. keep thermodynamics delayed until reversible support and detailed-balance behavior at $F_{\max} = 0$ are verified.

The most defensible statement is now:

In tested first-order regimes, topology and masks shape the rooted tree-sum basis used for ICL. Existing-data tree-difference multiplicity is predictive, but the first prospective exact-control contrast did not support it as a standalone causal knob. The best current pre-training signal is weak normal-fan/tree-count geometry, while the strongest mechanism evidence remains post-training branch/projection dependence.

[20] May 13 multi-base exact-control handoff

The next clean test has now been constructed. The file

`ICL/multibase_normal_fan_tree_count_library.py`

builds a multi-base exact-degree library whose purpose is to separate total rooted-tree abundance from task-aligned normal-fan/active-tree geometry. This is a targeted exact-control design, not a broad sweep.

The controls are:

$$N_n = 5, \quad m = 12, \quad N_c = 3, \quad D = 2, \quad p = 8, \quad d_{\text{rel}} = 88,$$

with full input coupling, exact multiplicity $M_\alpha = m$ for every input coordinate, and exact in/out degree sequence within each base graph. The design uses seven base graph families, 80 directed degree-preserving rewires per base, and paired selection arms:

arm	pairs	mean $ \Delta \text{ tree count} $	mean $ \Delta \text{ normal-fan score} $
fixed tree count, variable normal fan	11	0.909	4.353
variable tree count, matched normal fan	9	14.667	0.060

The resulting handoff contains 560 candidate topologies, 37 selected topology groups, 20 matched pairs, and 185 training tasks. The primary analysis plan is:

controls/base only,
tree count + base,
normal fan + base,
tree count + normal fan + base,
cross-root rank + tree count + normal fan + base.

The outcomes remain mean novel-class ICL, best-seed novel-class ICL, and seed standard deviation, using grouped LOO, held-out-base checks, matched-pair contrasts, and clustered bootstrap by base.

The local handoff paths are:

artifact	path
training plan	<code>ICL/results/next_phase_stats/multibase_normal_fan_tree_count_training_plan.md</code>
library report	<code>ICL/results/next_phase_stats/multibase_normal_fan_tree_count_separation_library.md</code>
library directory	<code>ICL/results/next_phase_stats/multibase_normal_fan_tree_count_n5_m12_N3_D2</code>
array script	<code>_array_meta/run_task.sh</code>

Operational status matters. The first Engaging submission, Slurm array 13858996, failed immediately because the generated script used system Python and did not load the Miniforge environment containing `numpy/torch`. The generator and concrete array script were patched to load `miniforge/25.11.0-0`. The corrected array was submitted as Slurm job 13902539; at the time of

Next exact-control library separates rooted-tree abundance from normal-fan geometry




Figure 16: Multi-base exact-control design. Arm A holds rooted-tree count nearly fixed while varying normal-fan score. Arm B varies rooted-tree count while matching normal-fan score. This is the next test for whether the weak normal-fan/tree-count signal is abundance, branch geometry, both, or neither.

this report update, the first 24 tasks had completed with exit code 0 and the remaining tasks were queued behind the array concurrency limit. No multi-base training result should be claimed until all 185 result files are collected.

The new current state is therefore:

The theory is exact at the tree-sum level; repaired gamma is analytic-toy valid but not predictive; same-root tree-difference overlap failed as a standalone causal knob; post-training branch/projection dependence remains strong; and the next live experiment is a multi-base exact-degree library designed to separate total rooted-tree count from normal-fan branch geometry.